

oBIX Tech Deep-Dive Document

Qingtao Cao [harry.cao@nextdc.com]

Last updated: 2014-11-05

Table of Contents

1. History Subsystem.....	3
1.1 Software Infrastructure.....	3
1.2 Filesystem Layout.....	4
1.3 XML DOM Hierarchy.....	4
1.4 Software Descriptors.....	5
1.5 Multi-thread Support.....	5
1.6 Performance of the Query Operation.....	5
2. Watch Subsystem.....	6
2.1 Software Infrastructure.....	6
2.2 XML DOM Hierarchy.....	8
2.3 Software Descriptors.....	8
2.4 Poll Threads and Poll Tasks.....	9
2.5 Extensible, Recyclable Bitmaps.....	9
3. Client Side Core Structures.....	10
3.1 Software Infrastructure.....	10
3.2 Software Descriptors.....	12
3.2.1 Comm_Stack.....	12
3.2.2 Connection and Http_Connection.....	12
3.2.3 Device and Http_Device.....	13
3.2.4 Listener and Http_Listener.....	13
3.3 CURL Handles.....	14
3.4 Race Condition on Listeners List.....	15
4. Modbus Gateway Adaptor.....	17
4.1 Hardware Connection.....	17

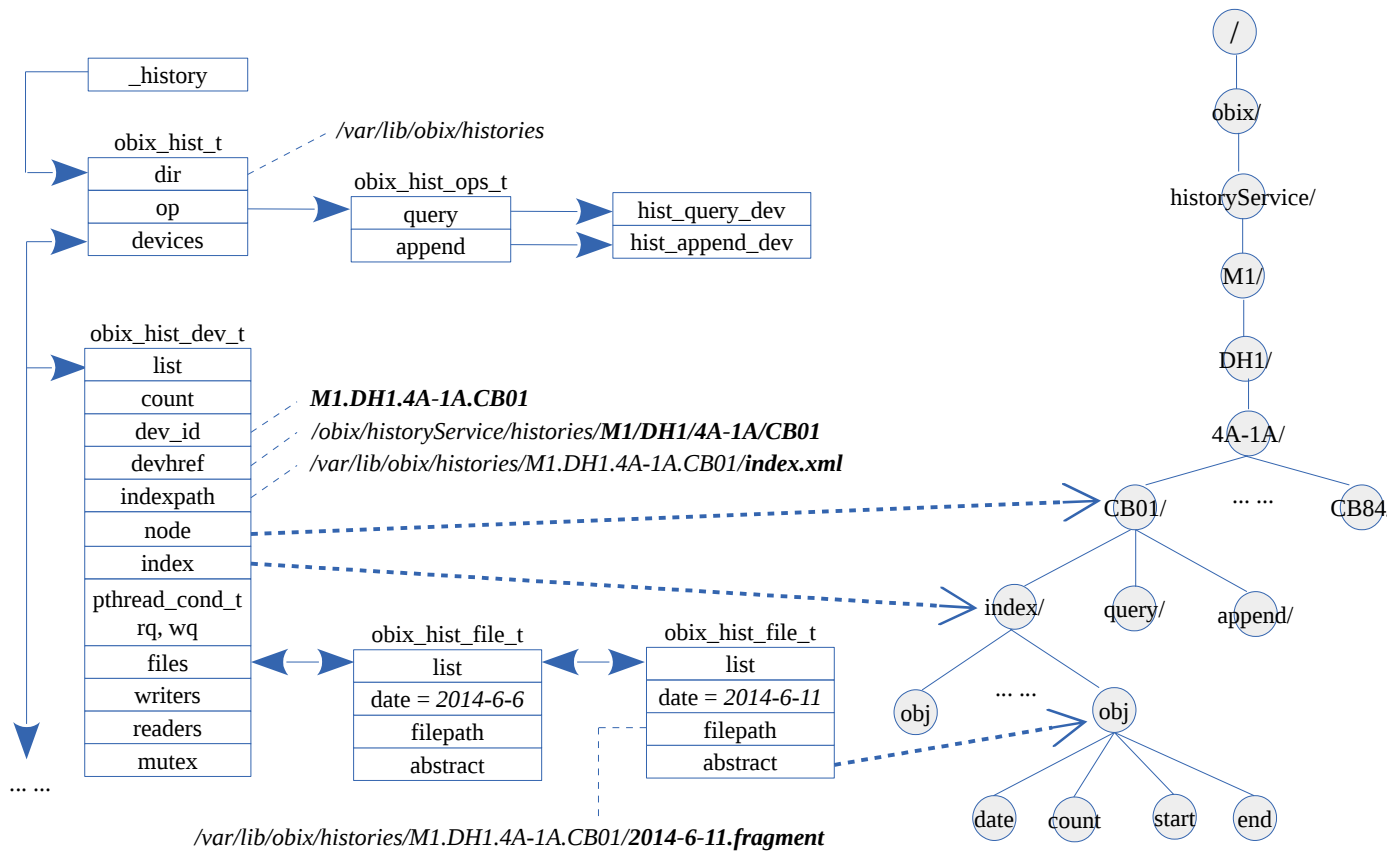
4.2 Interaction between Hardware and Software.....	19
4.3 oBIX Contracts for a BCM and a BM.....	20
4.4 Software Infrastructure.....	21
4.5 Producer-Consumer Model.....	22

Source Code:

`git@github.com:ONEDC/obix.git`

1. History Subsystem

1.1 Software Infrastructure



```

/var/lib/obix/histories/
├── M1.DH1.4A-1A.CB01
├── ...
├── 2014-06-06.fragment
├── 2014-06-11.fragment
└── index.xml

```

```

<obj/>
.....
<obj is="obix:HistoryRecord">
  <abstime name="timestamp" val="2014-06-11T07:40:05"/>
  <real name="kWh" val="0.000000"/>
  <real name="kW" val="0.000000"/>
  <real name="V" val="230.000000"/>
  <real name="PF" val="0.900000"/>
  <real name="I" val="0.000000"/>
</obj>

```

```

<?xml version="1.0" encoding="UTF-8"?>
<list name="index" href="index" of="obix:HistoryFileAbstract">
  <obj is="obix:HistoryFileAbstract">
    <date name="date" val="2014-06-06"/>
    <int name="count" val="577"/>
    <abstime name="start" val="2014-06-06T01:51:33"/>
    <abstime name="end" val="2014-06-06T06:39:02"/>
  </obj>
  <obj is="obix:HistoryFileAbstract">
    <date name="date" val="2014-06-11"/>
    <int name="count" val="810"/>
    <abstime name="start" val="2014-06-11T00:35:49"/>
    <abstime name="end" val="2014-06-11T07:40:05"/>
  </obj>
</list>

```

1.2 Filesystem Layout

The rightmost part of the above diagram shows the layout of history related folders and files on the hard-drive. oBIX clients can request the oBIX server to create a history facility for one particular device. On the hard-drive of the oBIX server, all such history facilities for different devices are organised under a particular folder.

As in this example, under the folder of `"/var/lib/obix/histories/"` there is a history facility named `"M1.DH1.4A-1A.CB01"`, which stands for the history facility for the Circuit Breaker 1 in a Branch Circuit Monitor named `"4A-1A"` in the Data Hall 1 of the M1 data centre.

This history facility contains one and only one index file in XML format and a number of raw history data files, each of which is nothing but a huge collection of XML objects representing history records. Since they don't have the required XML root elements, they are just XML fragments, that's why their file names are suffixed by `"fragment"`. BTW, each history data file is named after the date when they were generated.

For XML objects representing history records, each of them consists of a timestamp element describing its generation moment and a number of other elements as its payload. For example, the history record for a Circuit Breaker contains all the information about its power output at a particular moment, such as kWh, kW, Voltage, Power Factor and Current Intensity. As in this example, probably because no racks have been connected to this Circuit Breaker ever, both its kWh and Current Intensity values are zero.

The index file of a history facility contains abstract objects for all its fragment files. As in this example, they are for the raw history data files generated on 2014-6-6 and 2014-6-11 respectively. Aside from the generation date, the abstract object also shows the total number of records in a fragment file, the start and the end timestamp of the very first and the very last record in that file respectively.

Whenever a new history record is added to a history facility, it is actually appended to the very end of the **latest** fragment file with relevant abstract object updated accordingly. For instance, at least the count number is increased by 1 and the end timestamp is set equal to that of the newly appended history record.

BTW, it's worthwhile to mention that the history subsystem of the oBIX server enforces strict control on history records so that they are in strict timestamp ascending order. Therefore oBIX clients can't append a history record with timestamp older than that of the very end record in the latest fragment file.

Furthermore, if the new history record is on a brand new date, a new fragment file is created from scratch for that date with its abstract object inserted into the index file.

1.3 XML DOM Hierarchy

The middle part of the diagram is a segment of the global DOM tree related with the history subsystem. The history facility for the Circuit Breaker 1 mentioned earlier has its own XML sub-tree, which has two child nodes named `"query"` and `"append"`, representing the operations supported by this history facility, and the content of the index sub-tree is directly converted from relevant index file on the hard-drive.

It's important to note that for the sake of performance and efficiency, only index files are loaded into the global DOM tree at the oBIX server's start-up. Whereas history fragment files may contain hundreds of thousands of records with several GB data, they are **never** loaded into the global DOM tree.

1.4 Software Descriptors

As shown in the leftmost part of the above diagram, an `obix_hist_dev_t` structure is created for every existing history facility on the hard-drive. They are all organised into a "devices" list in the `obix_hist_t` structure for the entire history subsystem. This way, the history subsystem imposes no limitation on the number of history facilities on the oBIX server.

For every fragment file an `obix_hist_file_t` structure is created and they are all organised into the "files" list in the `obix_hist_dev_t` descriptor for the relevant history facility. This way, there is no limitation on the number of fragment files contained in one single history facility.

The `obix_hist_dev_t` descriptor describes a history facility's name, its parent href in the global DOM tree and the absolute path of its index file on the hard-drive, while the `obix_hist_file_t` descriptor contains the information about a fragment file's generation date and its absolute path in the filesystem. Both software descriptors contain pointers pointing to relevant XML nodes in the global DOM tree, which are highlighted by bold, dotted and blue lines with arrow in the above diagram.

Whenever a new history record is appended, both software descriptors and relevant XML nodes are updated altogether, and the content of the index sub-tree is also converted into a XML document and saved into relevant index file on the hard-drive to prevent potential data loss.

1.5 Multi-thread Support

In the first place, different history facilities are allowed to be accessed in parallel.

Moreover, in order to support concurrent access on one history facility from multiple oBIX clients who may query from and append to it simultaneously, the `obix_hist_dev_t` descriptor records its state information and employs POSIX pthread mutex and conditionals to synchronise among readers and writers.

In current implementation, the number of existing readers and the number of waiting writers are counted and multiple readers are permitted to read from one history facility in parallel. However, writers are excluded from each other and from any readers. Furthermore, no writers will suffer from starvation since no more readers will be allowed if there is any writers waiting for the completion of existing readers, in which case new readers will have to wait until each waiting writer has a chance to finish its task first.

1.6 Performance of the Query Operation

Having said earlier, a history facility may contain unlimited number of raw fragment files, each of which may further contain a vast number of records. So it's not unusual for one history facility to host several GB data, in this case the speed of the query operation could easily become a bottleneck of the entire history subsystem.

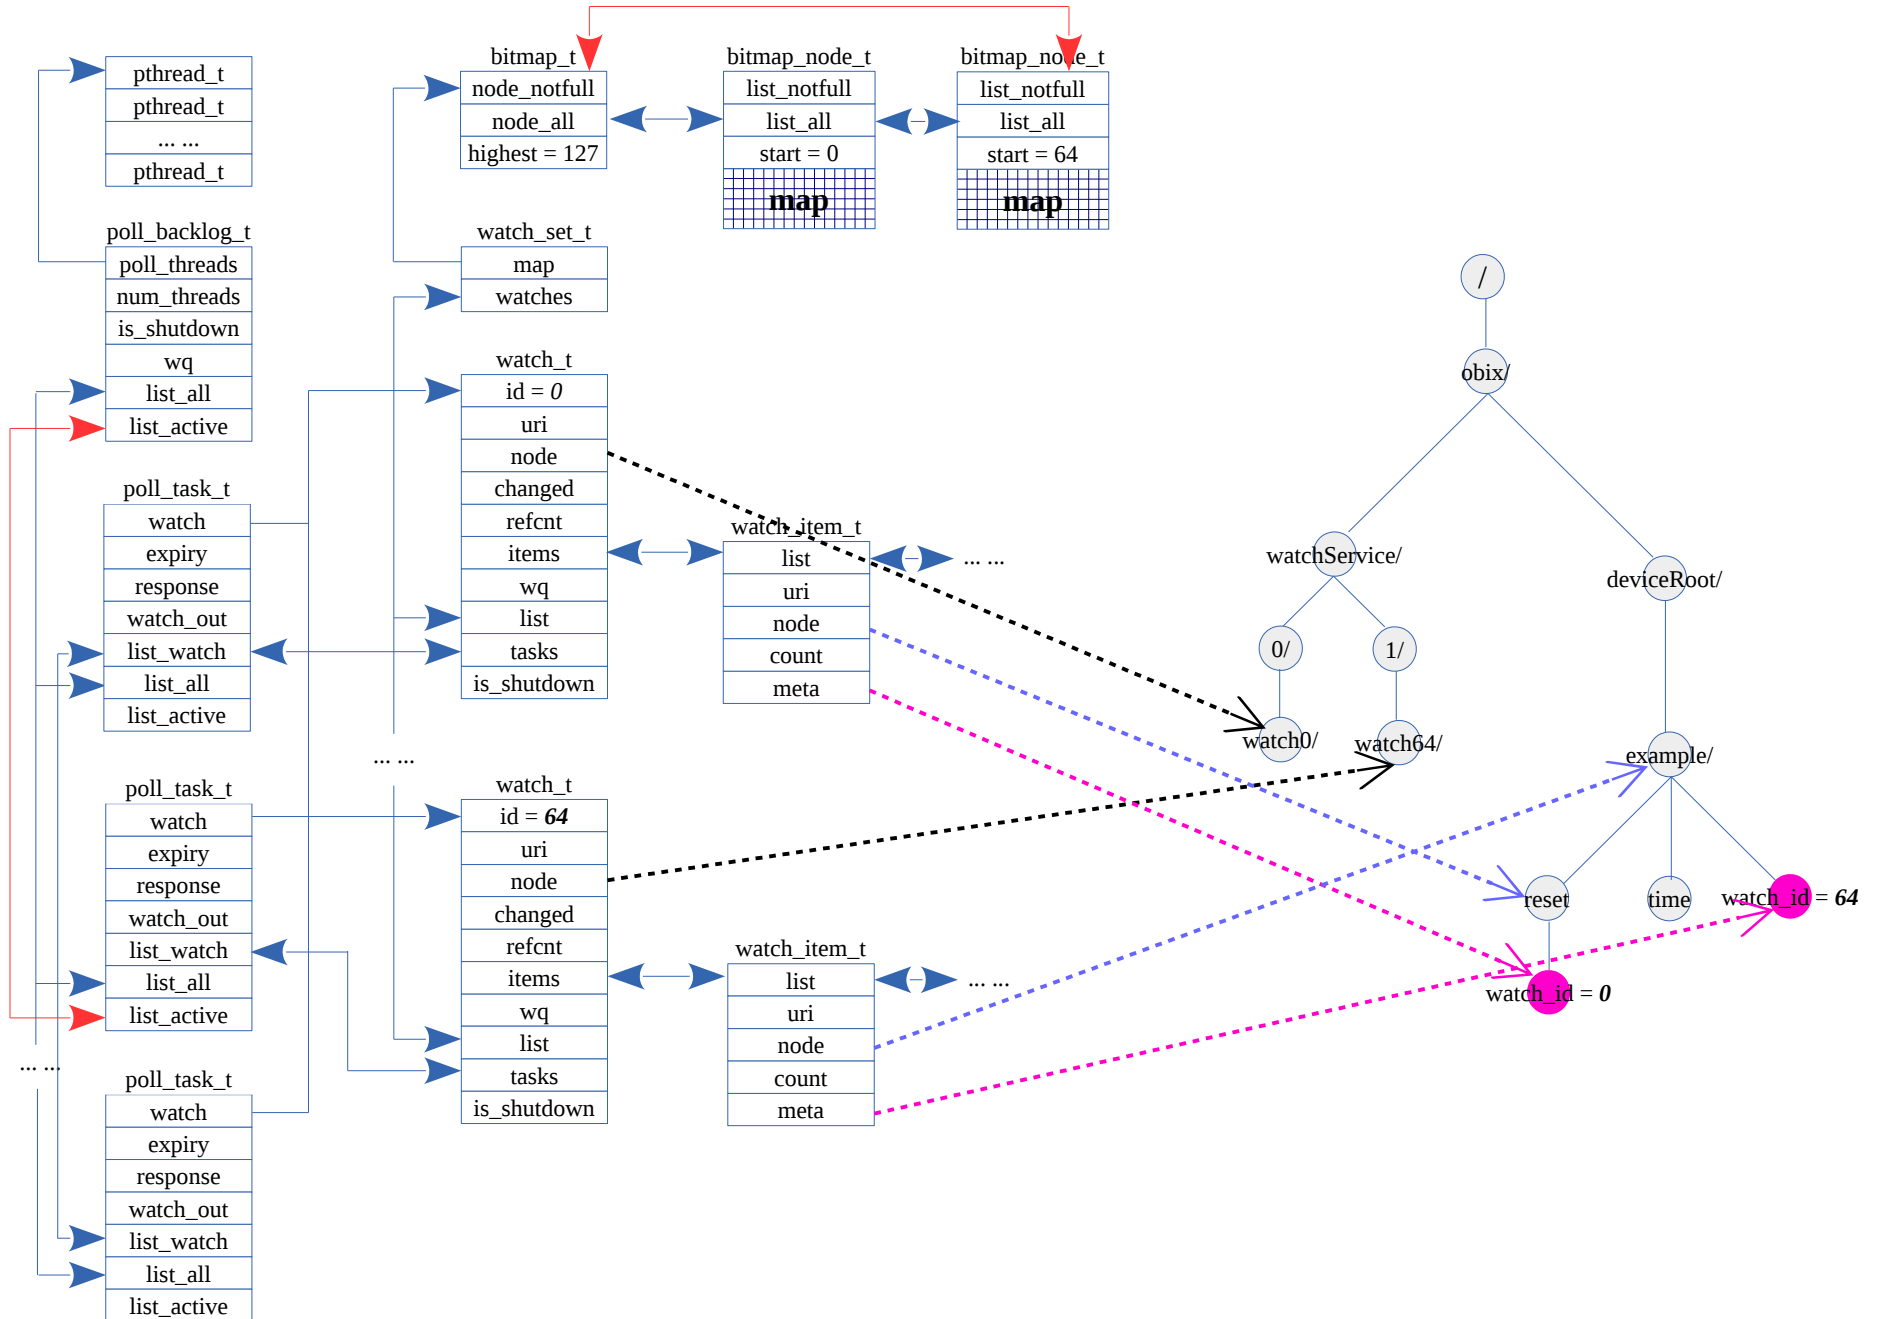
When an oBIX client is querying from a history facility, it needs to specify a limit on the number of records and a timestamp range. The query operation will only return satisfactory records from the specified history facility. Suppose a number of fragment files are involved in one query request, thanks to the fact that records in a fragment file are in strict timestamp ascending order and the fragment files are organised by their generation date, normally speaking, only the records in the first and the last fragment

files may need to be examined one after the other, whereas the content of all the rest fragment files in the middle could be directly returned back to relevant oBIX client.

Please refer to the source code for details about all those complicated mess of thoughts involved.

2. Watch Subsystem

2.1 Software Infrastructure



2.2 XML DOM Hierarchy

As shown by the right part of the above diagram, under the oBIX lobby there are two nodes for the WatchService and the deviceRoot respectively. The watchService href contains all XML contracts for every single watch objects already created on the oBIX server. Each watch object has a unique ID number and no more than 64 watch objects are grouped under one parent href. When there are tens of thousands of watch objects on one oBIX server and if they were simply organised as direct children of the watchService href, it would take ages to traverse the entire watchService sub-tree to look for one particular watch object with a specific ID number. That's why this 2-level hierarchy structure is adopted in order to promote performance.

The deviceRoot node hosts all XML contracts registered by different oBIX clients. As shown in the diagram there is an “example” device registered under it with two children nodes of reset and time. Suppose the “reset” node is currently monitored by a watch object with ID 0 and the entire example device is watched upon by another watch object with ID 64, a special type of meta element is installed under each of them to store the ID information of the relevant watch object. This way, whenever a XML node in the global DOM tree are changed, the oBIX server has knowledge about whether it is being monitored by any watch object and which watch object should be notified when changes occur.

It's also interesting to note what happens when two watch objects are monitoring the nodes in one same sub-tree at the same time. With the help of the meta elements installed in the global DOM tree, the oBIX server also has a clear idea of whether there is any watch object monitoring any **ancestor** nodes of the changed node. If this is the case, then the watch objects on the ancestors will be notified as well. For example, when the reset node is changed, both watch 0 and watch 64 are notified. However, only watch 64 is notified if the rest of the example device other than the reset node is changed.

BTW, if there are more than one watch objects monitoring one same XML node, multiple meta elements are created under it for each watch object respectively. And a meta element will be removed when relevant watch object no longer keeps an eye on that XML node.

2.3 Software Descriptors

The watch_set_t structure is the high-level descriptor of the entire watch subsystem, it has a “map” pointer pointing to a list of extensible bitmap nodes which provide ID numbers for new watch objects. When any watch object is deleted, its ID number is recycled and reused properly.

Whenever a brand-new watch object is created upon a request from an oBIX client, a watch_t descriptor is created and organised into the “watches” list in the watch_set_t descriptor. That's how the watch subsystem supports unlimited number of watches on the oBIX server.

For every XML node monitored by one watch object, a watch_item_t descriptor is created accordingly and they are organised into the “items” list in relevant watch_t descriptor. This way one watch object is able to monitor unlimited number of XML nodes in the global DOM tree.

Both the watch_t descriptor and the watch_item_t descriptor have pointers pointing to relevant XML nodes in the global DOM tree. In particular, the “node” pointer in the watch_t descriptor points to the watch object under the watchService href, while the “node” pointer in the watch_item_t descriptor points to the monitored XML node. These pointers are represented by bold, dotted, black or blue lines with arrow in the above diagram.

Furthermore, the `watch_item_t` descriptor also has a pointer named "meta" pointing to the special meta element installed under the monitored XML nodes in the global DOM tree. In the above diagram they are represented by bold, dotted and pink lines with arrow.

2.4 Poll Threads and Poll Tasks

The `poll_backlog_t` structure is the high-level descriptor of relevant polling threads and all pending poll tasks. Whenever an oBIX client would like to use a watch object to monitor some XML nodes in the global DOM tree, it firstly requests to have a watch object created, then tells it which XML node or nodes to monitor. Finally the client forks a thread blocking to listen to the notification sent back from the watch object. Accordingly, a `poll_task_t` descriptor is created on the oBIX server that contains all the information about which watch object is being polled upon, the maximal waiting time of relevant poll task and a pointer to the FCGI response structure which will carry the notification back to relevant oBIX client.

As a matter of fact, one `poll_task_t` structure can join as many as 3 different lists for different purposes.

In the first place, in order to support having one watch object shared among multiple oBIX clients, a `poll_task_t` descriptor is equipped with a "list_watch" field in order to join a "tasks" list in the relevant `watch_t` descriptor. When a watch object is being polled by multiple oBIX clients, relevant poll tasks are all organised in this queue. So whenever any watch item descriptor of this watch object detects any positive changes, all poll tasks in that queue are asserted so as to be further handled by a polling thread, so that in the very end, all oBIX clients sharing this watch object will get the same notification of changes.

In the second place, in order to enable the polling threads to take care of all polling tasks in an efficient manner, each `poll_task_t` descriptor can join two queues at the same time. In particular, upon creation a `poll_task_t` descriptor simply takes part in one universal "list_all" queue containing all poll tasks and its position in that queue is solely decided by its expiry date. This universal queue is organised according to the ascending order of the expiry date of each poll task, so that the polling threads are able to identify and handle those expired poll tasks just at the beginning of that queue effectively and efficiently.

Moreover, one `poll_task_t` descriptor will also join a separate, "list_active" queue right after relevant watch object is triggered by a change event, so that it is placed immediately on the radar of the polling threads, which will race to grab one poll task from the queue and handle it. It's worthwhile to mention that the polling threads always handle the active queue first, then the expired tasks at the beginning of the universal queue, then fall back asleep again until waken up by another change event or the first poll task(s) in the universal queue becomes expired.

The watch subsystem employs POSIX pthread mutex to eliminate potential race conditions among polling threads, utilises pthread conditionals to synchronise among polling threads and the oBIX server thread, and also manipulates a reference count to ensure a `watch_t` descriptor won't get destroyed until all relevant pending poll tasks have been properly handled.

2.5 Extensible, Recyclable Bitmaps

Previously a simple integer was used to count the ID number for watch objects, it would overflow in the future when a large number of watch objects were created and then deleted repeatedly, consequently, newly created watch objects may have same ID number as existing ones which is not acceptable at all since it is critical for a watch object

to be assigned a unique ID number so as to be used by oBIX clients properly. To address this challenge a list of bitmap nodes are adopted, each of which covers 64 numbers from its starting number and they are organised in a “node_all” queue in the ascending order of their starting number.

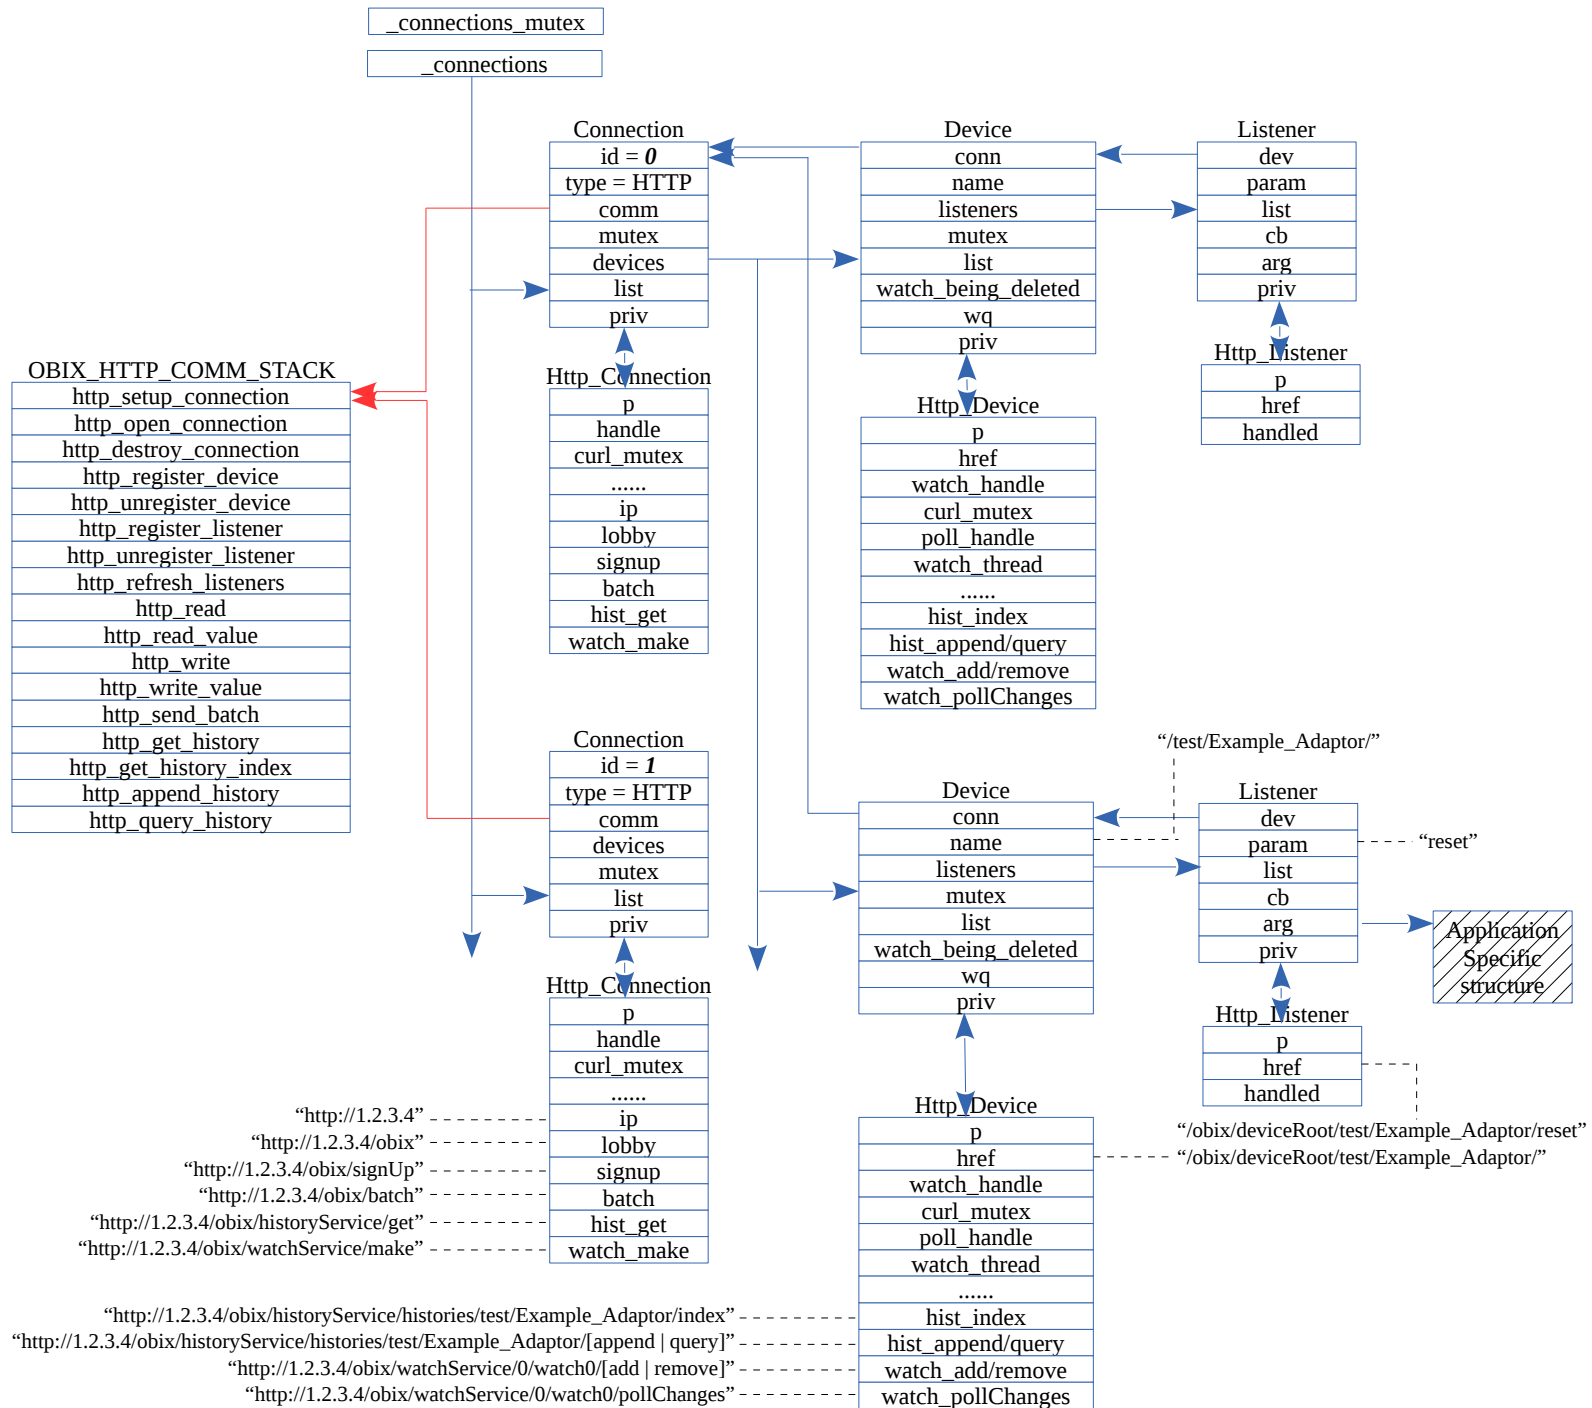
If a number is picked up for a new watch object, a bit in relevant bitmap node is asserted to 1, whereas whenever a watch object is deleted, its ID number is returned, resulting in relevant bit in relevant bitmap node is reset to 0.

If any bitmap node has more than one unused number, it also joins a separate “node_notfull” queue which is also organised in strict ascending order of each bitmap node's starting number. This way, the ***smallest*** unused number is always returned and used.

Lastly, if a bitmap node has run out of available bits, it will be removed from the “node_notfull” queue. Of course, it will re-join this queue whenever any of its bits are returned and reset to 0 so as to be reused later.

3. Client Side Core Structures

3.1 Software Infrastructure



3.2 Software Descriptors

The high-level oBIX software descriptors such as Connection, Device and Listener are low-level or binding neutral, which aim to describe a connection with an oBIX server, a device contract registered on it and a child node of that device monitored by a watch object on the server respectively. Meanwhile, they rely on low-level, HTTP counterparts such as Http_Connection, Http_Device and Http_Listener to host respective HTTP specific details on different identities.

3.2.1 Comm_Stack

The set of function pointers in the Comm_Stack structure provide an interface between high-level, oBIX APIs and low-level, binding specific implementation. The oBIX project implements only HTTP bindings for the time being and relevant methods are organised within an OBIX_HTTP_COMM_STACK structure. If SOAP bindings were to develop in the future, they should be arranged in a similar structure and a connection could selectively use one of supported bindings according to its relevant configuration settings.

Generally speaking, oBIX APIs manipulate high-level data structures and further invoke low-level, HTTP APIs to manipulate low-level data structures to fulfill particular operations. However, this is not necessarily always true. For example, obix_get_history_ts() simply invokes another high-level API of obix_query_history() with a special historyFilter contract to get timestamp strings of the first and last records from relevant history facility directly, without involving extra low-level operations.

3.2.2 Connection and Http_Connection

oBIX client applications can setup connections with more than one oBIX server, for instance, one oBIX Adaptor that needs to push hardware data to multiple oBIX servers simultaneously, or the “Synchronise Adaptor” which normally picks up particular information from a low-level oBIX server and relay them onward to another one at a higher level. To this end, Connections are organised into a global linked list to support no limitation on the number of them and establishment or deletion of connections at run-time, although oBIX clients practically setup their connections at start-up and use them throughout their life cycle.

Currently oBIX clients utilise a XML configuration file to specify details of connections with different oBIX servers, such as their unique ID number, IP address, type information etc, they are stored into connection related structures. Having said this, for the time being only HTTP binding is implemented and therefore all Connections' “comm” pointers point to the same OBIX_HTTP_COMM_STACK structure.

The Http_Connection structure describes the HTTP specific details of a connection, such as relevant oBIX server's IP address, URIs for all public facilities on that oBIX server and the particular CURL handle created and used exclusively for this connection. All URIs are prefixed by the server's IP address for the convenience of libcurl APIs. Please refer to the following section about CURL handle for pitfalls when manipulating them.

3.2.3 Device and Http_Device

The Device structure represents the high-level information of a device (or a XML object or an oBIX contract) registered to a particular oBIX server. Devices registered through one Connection are organised in its “devices” linked list, which is protected by a mutex to enable dynamic Device insertion and deletion. Each Device uses a “conn” pointer pointing back to the parent Connection.

Each Device has a unique name on its connection with an oBIX server as specified by relevant configuration settings. However, devices on different connections with the same server may happen to have duplicated names, in order to ensure devices are recognisable by their names on the server, their names are always prefixed by another “parent_href” configuration settings which specifies the root location of the registered device contract under the global device lobby of the oBIX server. In theory IT administrator must ensure the uniqueness of the “parent_href” settings of different oBIX clients across the whole server.

The Http_Device structure describes the HTTP specific details of a device, such as its href on relevant oBIX server, URIs for various facilities supported by this device, the worker thread dedicated for watch.PollChanges requests and CURL handles used for raising blocking and non-blocking requests respectively.

3.2.4 Listener and Http_Listener

The Listener structure is the client side equivalent to the server side watchItem structure, containing high-level information of a monitored child node of a particular Device, such as its relative href on the Device, the callback function invoked to claim changes and relevant parameter. Generally speaking the parameter can point to application specific data structure used by the callback function to handle changes on the monitored child node.

Generally speaking, multiple Listeners can be registered to monitor different child nodes of a Device independently, they are organised in the Device's “listeners” linked list, which is protected by a mutex to enable dynamic insertion and deletion. Each Listener uses a “dev” pointer pointing back to the parent Device. The Device's Listeners list is traversed once its worker thread receives a change notification (the watchOut contract to be exact, which may contain a list of updated nodes on the Device) from the oBIX server and those concerned callback functions are invoked one after the other.

It's worthwhile to note that the client side Device specific worker thread, the CURL handles related with watch relevant requests and the server side watch object are all created in a *lazy* mode, that is, they are established and destroyed along with the first Listener and the last Listener respectively. So the needed infrastructure is set up only when they are really needed and applications that never make use of them won't waste any memory at all.

Last but not least, a special use case is to set “param” argument equal to “/” which will have the entire device monitored instead of any child node, since double slashes are regarded as one when assembling the href of the monitored node therefore the href of the whole device is watched upon.

3.3 *CURL Handles*

Firstly and most importantly, CURL handles are ***not*** thread-safe therefore multi-thread applications must synchronise accesses on them. Each thread can either have its own exclusive CURL handle or contend for a mutex before accessing a shared one.

For sake of performance applications are encouraged to use multiple CURL handles when necessary, especially when it comes to blocking requests such as `watch.PollChanges` when requesters are likely to be blocked by libcurl API until a notification is returned back from the oBIX server, in which case each requester should independently manipulate their own CURL handle in order to receive each of their respective notifications in a timely manner.

To this end, a `Http_Device` structure hosts a “poll_handle” CURL handle exclusively used by its own “watch_thread” worker thread to listen for any changes occurred on the associated parent Device. Furthermore, a `Http_Device` structure also hosts another CURL handle named “watch_handle” so that other non-blocking watch requests can be raised in parallel with any blocking `watch.PollChanges` requests on the same Device.

Since Listeners are added and deleted at run-time, the `Http_Device` structure uses a mutex “curl_mutex” to serialise accesses to the “watch_handle”. Whereas the “poll_handle” is solely used by the worker thread so no such protection is required.

Lastly, many oBIX APIs accept the address of a user provided CURL handle as the first parameter. If it is NULL, then Connection specific or Device specific default CURL handle is fallen back on. Otherwise user applications must synchronise their usage on their own CURL handle.

The CURL handles used for different types of requests are summarised in below table:

Request Type		CURL Handle Used
GET		User defined or Connection default
PUT		User defined or Connection default
POST	signup	Connection default
	batch	User defined or Connection default
	history.Get	User defined or Connection default
	history.Append	User defined or Connection default
	history.Query	User defined or Connection default
	watch.Make	Device's "watch_handle"
	watch.Add	Device's "watch_handle"
	watch.Remove	Device's "watch_handle"
	watch.PollRefresh	Device's "watch_handle"
	watch.PollChanges	Device's "poll_handle"
Listeners' callbacks		Device's "poll_handle" (*)

As we can see, the default CURL handle of the entire Connection is good for single-thread applications, whereas multi-thread applications must manipulate their own CURL handles to avoid racing against the Connection default one which will be a performance bottleneck, especially when it comes to time-consuming requests such as history.Query when massive amount of data may be transmitted over a slow network.

Note (*): once the Device specific worker thread has received the watchOut notification from oBIX server, it traverses the list of Listeners and invokes a Listener's callback function with its own CURL handle if it is concerned about current changes.

3.4 Race Condition on Listeners List

Attentions must be paid when accessing a Device's Listeners list, following threads are racing against each other on it:

- 1) application threads to create a new Listener (through the Device's "watch_handle");
- 2) application threads to remove an existing Listener (through the Device's "watch_handle");
- 3) the Device's specific worker thread polling for changes on the device (through the Device's "poll_handle").

In the first place, a mutex needs to be held in scenario #1 and another mutex is also needed to ensure synchronised access to the watch_handle. The former mutex must be held until the communication with the oBIX server is finished even if it may take a relatively long time. Come to think of it, if a thread is blocked waiting for responses from a busy oBIX server or through a slow connection, there won't be extra benefit to have other threads raising similar requests.

In the second place, the Device's worker thread will release the mutex before invoking a Listener's callback function, which is supposed to access application specific structures and be time-consuming. As a result, the Listeners list could be modified during the invocation of a callback function, therefore the worker thread must restart from the beginning of the list and skip over already handled ones. To this end, the Http_Listener structure hosts an indicator about whether it has been handled by this round of execution of the worker thread.

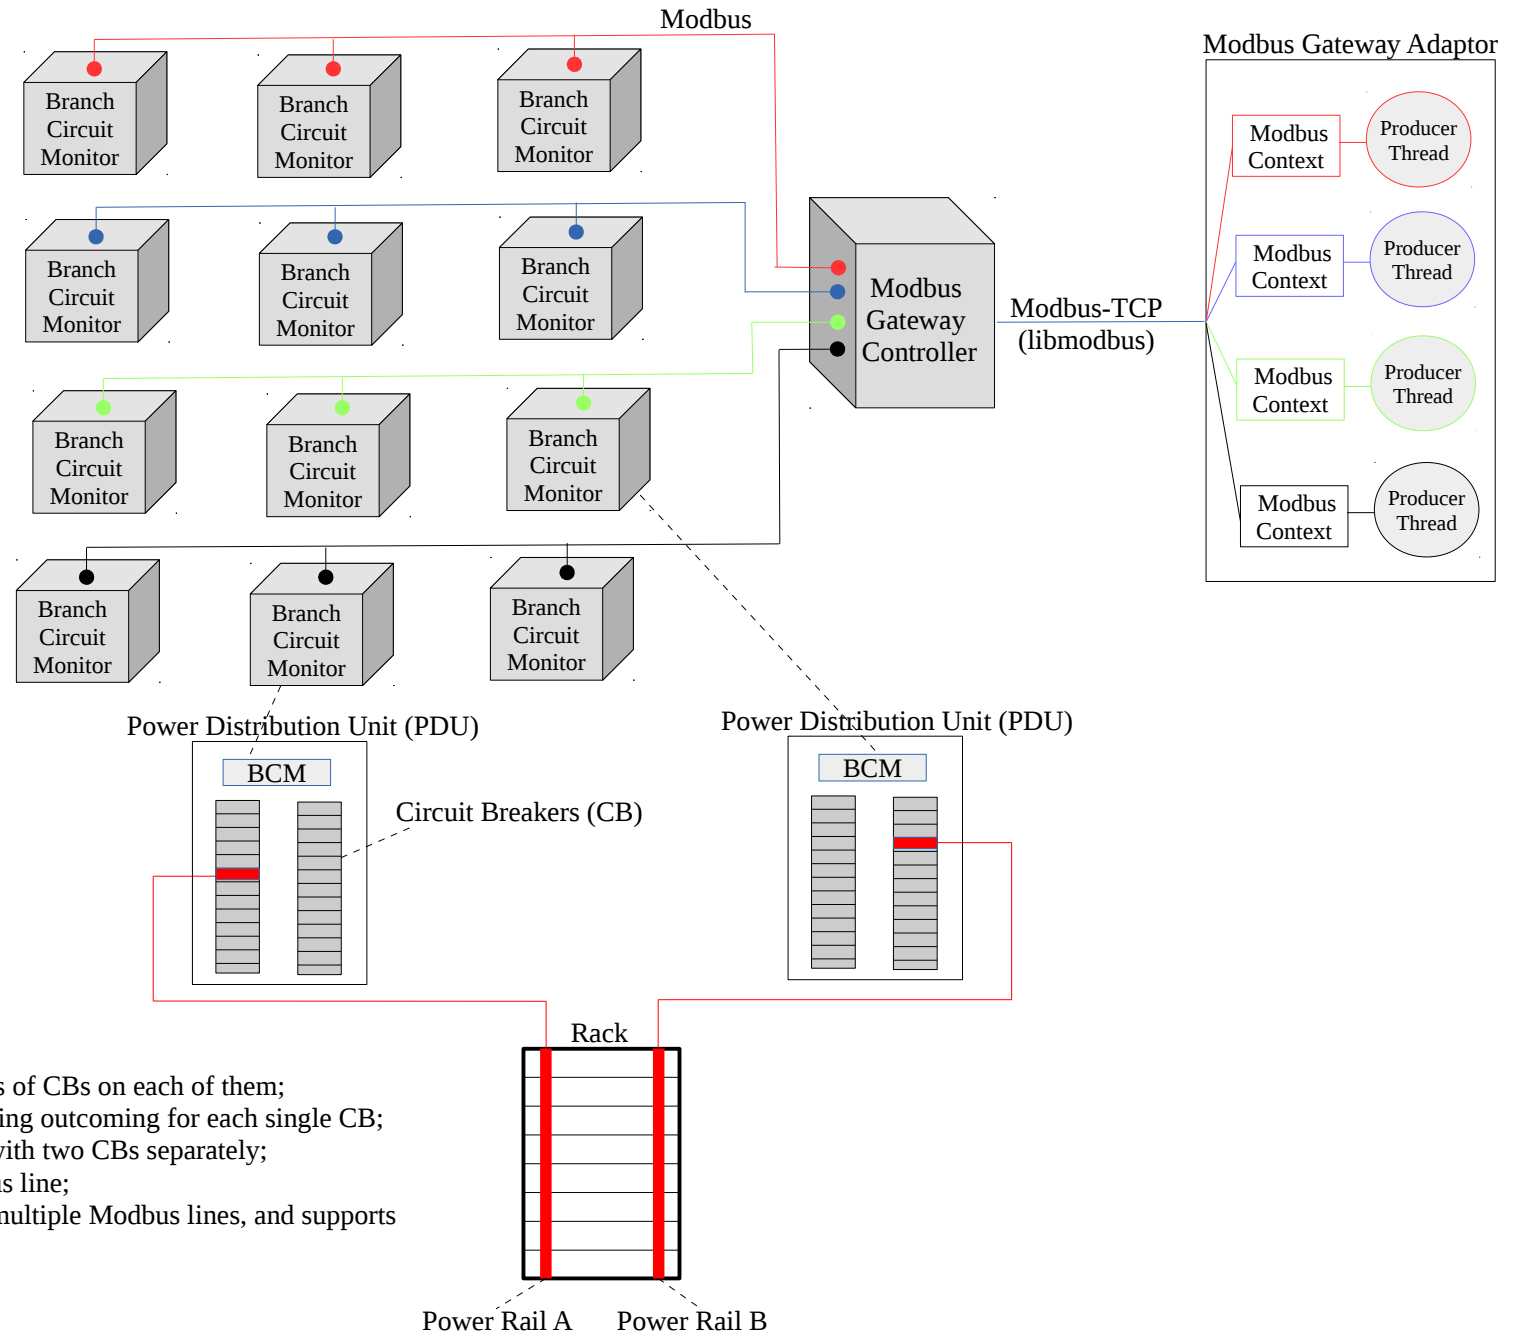
In the third place, if the to-be-deleted Listener is the last one in the queue, the client application will also request oBIX server to delete relevant watch object and then destroy relevant worker thread and Device specific CURL handles on the client side. In such case no worker thread should ever be running at back-end when relevant API returns. To this end, the deletion thread in scenario #2 needs to wait for the completion of the running worker thread before terminating it, which in turn requires the Device's mutex protecting the Listeners queue must be released by the deletion thread in order to avoid a deadlock with the worker thread who also competes on the same mutex.

However, above method invites a race condition. Once the Device's mutex is released by the deletion thread, another application thread in scenario #1 may be scheduled in immediately which tries to add a new Listener into the queue (and a new watchItem to relevant watch object on the server side). However, the entire watch object and relevant structures will be deleted shortly after when the mutex is released and the deletion thread is resumed.

To further address this race condition, a flag indicating whether the watch object on the server side is under deletion and a conditional are adopted. Any thread in scenario #1 requesting to create a new watch object would have to wait for the completion of the deletion of an existing one.


4. Modbus Gateway Adaptor

4.1 Hardware Connection



1. One PDU hosts two panels, with dozens of CBs on each of them;
2. One BCM embeds inside a PDU, metering outcoming for each single CB;
3. The two power rails of a rack connect with two CBs separately;
4. BCMs are daisy-chained on one Modbus line;
5. One Modbus gateway controller hosts multiple Modbus lines, and supports concurrent access to each of them.

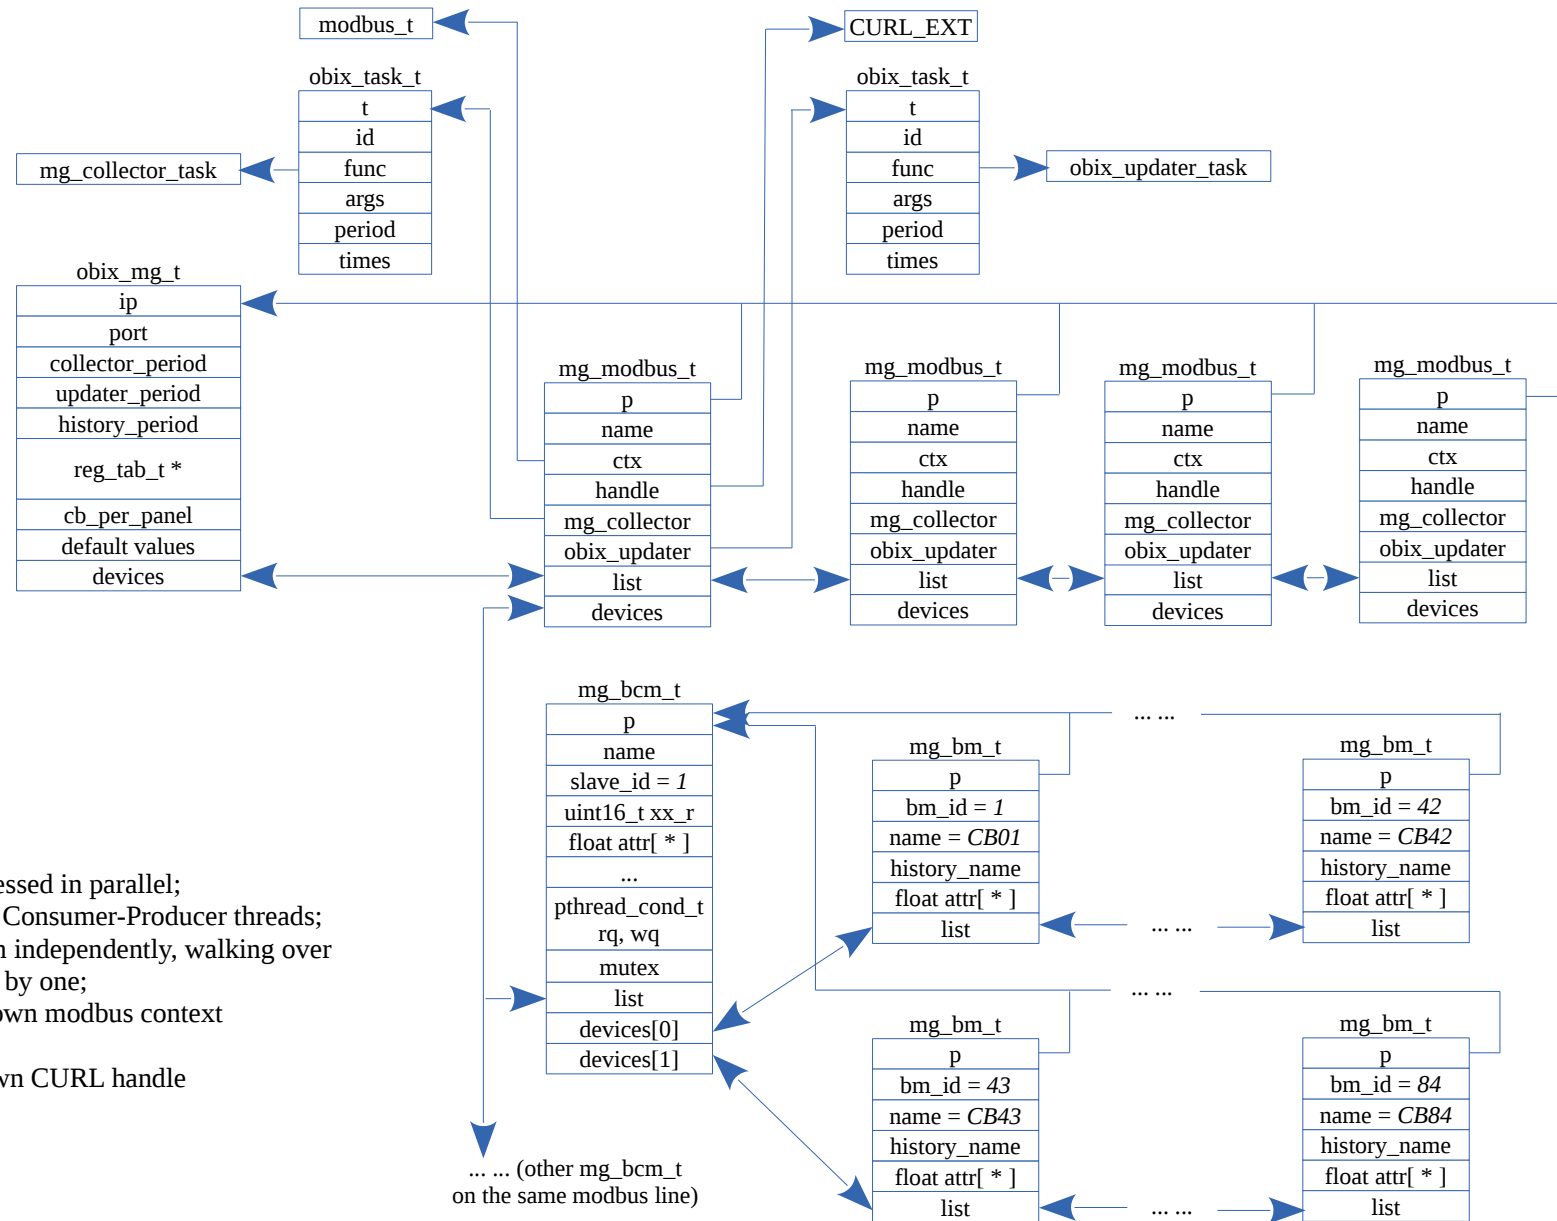
4.2 Interaction between Hardware and Software



4.3 oBIX Contracts for a BCM and a BM

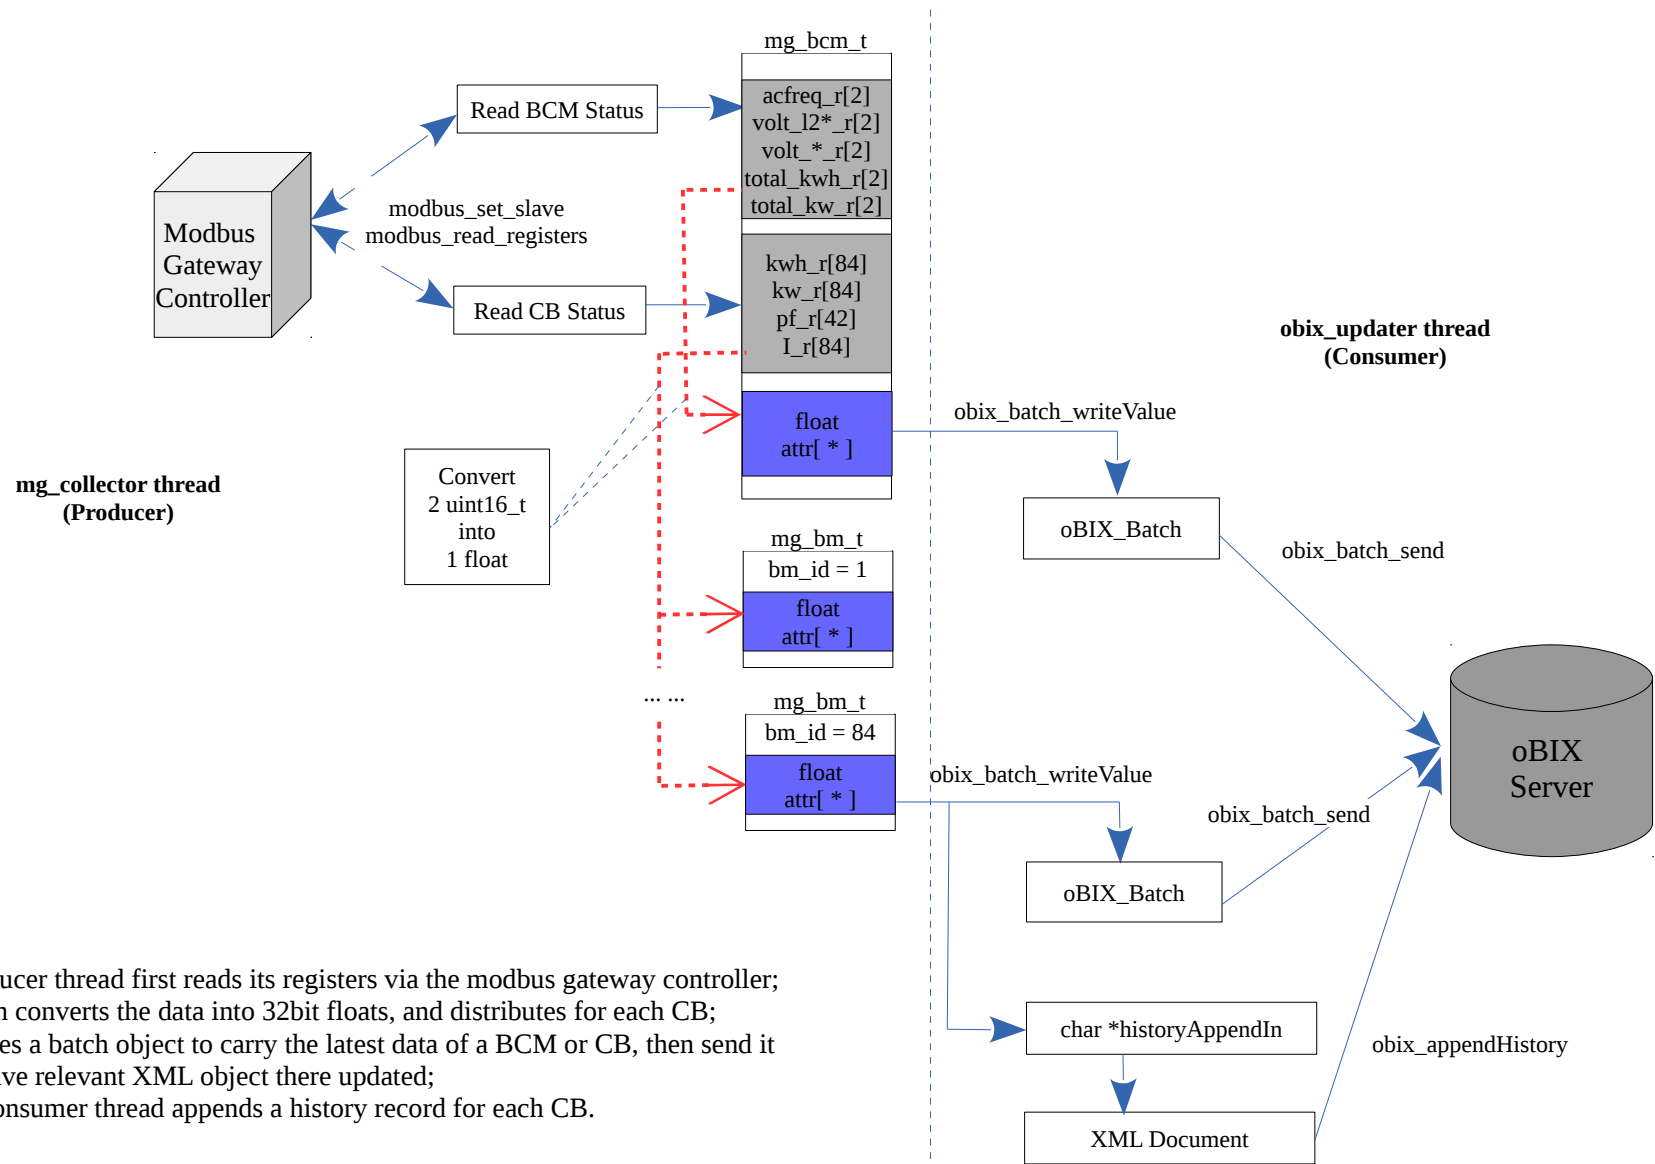
```
<obj name="4A-1A" href="/obix/deviceRoot/M1/DH4/4A-1A" is="nextdc:VerisBCM">
  <int name="SlaveID" href="SlaveID" val="1"/>
  <int name="SerialNumber" href="SerialNumber" val="0x4e342ef9" writable="true"/>
  <int name="Firmware" href="Firmware" val="0x03ed03f4" writable="true"/>
  <int name="Model" href="Model" val="15172" writable="true"/>
  <int name="CTConfig" href="CTConfig" val="2" writable="true"/>
  <str name="Location" href="Location" val="AUDM1DH4 PDU-4A-1A Panel #1" writable="true"/>
  <real name="ACFreq" href="ACFreq" val="50.000000" writable="true"/>
  <real name="VoltL-N" href="VoltL-N" val="242.080078" writable="true"/>
  <real name="VoltL-L" href="VoltL-L" val="418.952057" writable="true"/>
  <real name="VoltA" href="VoltA" val="240.157227" writable="true"/>
  <real name="VoltB" href="VoltB" val="243.659668" writable="true"/>
  <real name="VoltC" href="VoltC" val="242.563507" writable="true"/>
  <real name="kWh" href="kWh" val="218.000000" writable="true"/>
  <real name="kW" href="kW" val="0.000000" writable="true"/>
  <real name="CurrentAverage" href="CurrentAverage" val="0.000000" writable="true"/>
  <abstime name="LastUpdated" href="LastUpdated" val="2014-05-19T01:19:24Z" writable="true"/>
  <bool name="Online" href="OnLine" val="true" writable="true"/>
  <list name="Meters" href="Meters" of="nextdc:Meter">
    <obj name="CB01" href="CB01" is="nextdc:Meter">
      <real name="kWh" href="kWh" val="25.444157" writable="true"/>
      <real name="kW" href="kW" val="0.000000" writable="true"/>
      <real name="V" href="V" val="240.157227" writable="true"/>
      <real name="PF" href="PF" val="0.900000" writable="true"/>
      <real name="I" href="I" val="0.000000" writable="true"/>
    </obj>
    .....
    <obj name="CB84" href="CB84" is="nextdc:Meter">
      <real name="kWh" href="kWh" val="50.943935" writable="true"/>
      <real name="kW" href="kW" val="0.000000" writable="true"/>
      <real name="V" href="V" val="243.659668" writable="true"/>
      <real name="PF" href="PF" val="0.900000" writable="true"/>
      <real name="I" href="I" val="0.000000" writable="true"/>
    </obj>
  </list>
</obj>
```

4.4 Software Infrastructure



1. Different modbus lines are accessed in parallel;
2. Each modbus line has a pair of Consumer-Producer threads;
3. Consumer-Producer threads run independently, walking over BCMs on one modbus line one by one;
4. Each Consumer thread has its own modbus context (which is not thread safe);
5. Each Producer thread has its own CURL handle

4.5 Producer-Consumer Model



1. For each BCM, the Producer thread first reads its registers via the modbus gateway controller;
2. The Producer thread then converts the data into 32bit floats, and distributes for each CB;
3. The Consumer thread uses a batch object to carry the latest data of a BCM or CB, then send it to the oBIX server to have relevant XML object there updated;
4. At specified pace, the Consumer thread appends a history record for each CB.